

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/286933843>

A Survey on Working Principle and Application of Hadoop

Article · November 2015

CITATION

1

READS

186

5 authors, including:



[Surajit Mohanty](#)

Dhaneswar Rath Institute of Engineering and Management Studies

5 PUBLICATIONS 1 CITATION

SEE PROFILE

A Survey on Working Principle and Application of Hadoop

Surajit Mohanty¹, Girish Das², Himanshu Suman³, Piyush Maharana⁴, Raman Ratnakar⁵

¹⁻²Asst. professor, ³⁻⁵B_Tech Scholar

¹⁻⁵Department of Computer Science Engineering, DRIEMS, Cuttack, India

¹mohanty.surajit@gmail.com

²girish.das@driems.ac.in

³himanshu.suman90@gmail.com

⁴piyush.maharana@gmail.com

⁵ratnakar.raman@gmail.com

Abstract- Hadoop is a main buzz phrase and new curve for IT today. Big data is driven data with high velocity, volume, variety, veracity and value. It comes from different sources like mobile devices, internet, social media, sensors etc. Traditional data processing and analysis of structured data using RDBMS and data warehousing no longer satisfy the challenges of Big Data. Due to the high velocity and volume of big data, the effective option is to store the big data through Hadoop, because it has capability to store and process massive amount of big data. Hadoop offers the big data implementation in small and medium sized businesses. This paper presents the working with Hadoop and its implementation in various sectors that include healthcare, networking security, market and business, sports, education system, gaming and telecommunications.

Keywords- Big data, Hdfs, Streaming, MapReduce, Combiner

I. INTRODUCTION

Big Data is a collection of large or complex data sets that cannot be processed using traditional computing techniques. Challenges include analysis, capture, datacuration, search, sharing, storage, transfer, visualization and information privacy. It is not a single technique or a tool; rather it involves many areas of business and technology. We live in the data age.

II. HISTORY

Hadoop was created by Doug Cutting, the creator of Apache Lucene, which is a widely used text search library.[3] Hadoop find its roots to Apache Nutch, an open source web search engine, itself a part of the Lucene project. Nutch was started in 2002, and a working crawler and search system quickly emerged. However, its creators realized that their architecture won't be able to scale the billions of pages on the Web. Help was at hand with the publication of GFS in 2003 which was being used in production at Google. In 2004, Google published the paper that introduced MapReduce to the world. Early in 2005, the Nutch developers had a working MapReduce implementation in Nutch, and by the middle of that year all the major Nutch algorithms had been ported to run using MapReduce and NDFS. In February 2006 they left Nutch to form an independent subproject of Lucene called Hadoop. At around the same time, Doug Cutting joined Yahoo!, which provided a dedicated team and the resources to turn Hadoop into a system that would run at web scale. Hadoop was demonstrated in February 2008 when Yahoo! announced that its production search index was being generated by a 10,000-core Hadoop cluster. Cutting named it

after his son's toy elephant. It was originally developed to support distribution for the Nutch search engine project.

III. WORKING PRINCIPLES OF HADOOP

Hadoop is an Apache open source framework written in java that allows distributed processing of large datasets across clusters of computers using simple programming models [4]. The Hadoop framework application works in an environment that provides distributed storage and computation across clusters of computers. It is designed to scale up from single server to thousands of machines, each offering local computation and storage. Hadoop consists of the Hadoop Common package, which provides file system and OS level abstractions, a MapReduce engine (either MapReduce/MR1 or YARN/MR2) and the Hadoop Distributed File System (HDFS). The Hadoop Common package contains the necessary Java ARchive (JAR) files and scripts needed to start Hadoop which provides source code, documentation, and a contribution section that includes projects from the Hadoop Community.

A small Hadoop cluster includes a single master and multiple worker (slave) nodes. The master node consists of a job tracker and name node. A slave or worker node acts as both a data node and task tracker, though it is possible to have data-only worker nodes and compute-only worker nodes. These are normally used only in nonstandard applications. Hadoop requires Java Runtime Environment (JRE) 1.6 or higher. The standard startup and shutdown scripts require that Secure Shell (ssh) be set up between nodes in the cluster. In a larger cluster, the HDFS is managed through a dedicated name node server to host the file system index, and a secondary name node that can generate snapshots of the name node's memory structures, thus preventing file-system corruption and reducing loss of data. Similarly, a standalone job tracker server can manage job scheduling. In clusters where the Hadoop MapReduce engine is deployed against an alternate file system, the name node, secondary name node, and data node architecture of HDFS are replaced by the file-system-specific equivalents.

At its core, Hadoop has two major layers namely: Storage layer (Hadoop Distributed File System), and Processing/Computation layer (MapReduce).

A. Hadoop Distributed File System (HDFS)

An HDFS cluster has two types of nodes operating in a master-slave pattern: a *Name Node* (the master) and a number of *Data Nodes* (slave). The name node manages the file system namespace. It maintains the file system tree and the metadata for all the files and directories in the tree. This information is stored persistently on the local disk in the form of two files: the namespace image and the edit log. The name node also knows the data nodes on which all the blocks for a given file are located; however, it does not store block locations persistently, because this information is reconstructed from data nodes when the system starts.

The first way is to back up the files that make up the persistent state of the file system metadata. Hadoop can be configured so that the name node writes its persistent state to multiple file systems. These writes are synchronous and atomic.

It is also possible to run a *secondary name node*, which despite its name does not act as a name node. Its main role is to periodically merge the namespace image with the edit log to prevent the edit log from becoming too large. The secondary name node usually runs on a separate physical machine because it requires plenty of CPU and as much memory as the name node to perform the merge. It keeps a copy of the merged namespace image, which can be used in the event of the name node failing. However, the state of the secondary name node lags that of the primary, so in the event of total failure of the primary, data loss is almost certain. The usual course of action in this case is to copy the name node's metadata files that are on NFS to the secondary and run it as the new primary.

Block Caching: Normally a data node reads blocks from disk, but for frequently accessed files the blocks may be explicitly cached in the data node's memory, in an off-heap *block cache*. By default, a block is cached in only one data node's memory, although the number is configurable on a per-file basis. Job schedulers (for MapReduce, Spark, and other frameworks) can take advantage of cached blocks by running tasks on the data node where a block is cached, for increased read performance.

HDFS Federation: The name node keeps a reference to every file and block in the file system in memory, which means that on very large clusters with many files, memory becomes the limiting factor for scaling. HDFS federation, introduced in the 2.x release series, allows a cluster to scale by adding name nodes, each of which manages a portion of the file system namespace. Under federation, each name node manages a *namespace volume*, which is made up of the metadata for the namespace, and a *block pool* containing all the blocks for the files in the namespace. Namespace volumes are independent of each other, which means name nodes do not communicate with one another, and furthermore the failure of one name node does not affect the availability of the namespaces managed by other name nodes. Block pool storage is not

partitioned, however, so data nodes register with each name node in the cluster and store blocks from multiple block pools. To access a federated HDFS cluster, clients use client-side mount tables to map file paths to name nodes.

HDFS High Availability: The combination of replicating name node metadata on multiple file systems and using the secondary name node to create checkpoints protects against data loss, but it does not provide high availability of the file system. The name node is still a single point of failure (SPOF). If it did fail, all clients including MapReduce jobs would be unable to read, write, or list files, because the name node is the sole repository of the metadata and the file-to-block mapping. In such an event, the whole Hadoop system would effectively be out of service until a new name node could be brought online. To recover from a failed name node in this situation, an administrator starts a new primary name node with one of the file system metadata replicas and configures data nodes and clients to use this new name node. The new name node is not able to serve requests until it has

- I. loaded its namespace image into memory,
- II. replayed its edit log, and
- III. Received enough block reports from the data nodes to leave safe mode.

On large clusters with many files and blocks, the time it takes for a name node to start from cold can be 30 minutes or more. The long recovery time is a problem for routine maintenance, too. In fact, because unexpected failure of the name node is so rare, the case for planned downtime is actually more important in practice.

Hadoop 2 remedied this situation by adding support for HDFS high availability (HA). In this implementation, there is a pair of name nodes in an active-standby configuration. In the event of the failure of the active name node, the standby takes over its duties to continue servicing client requests without a significant interruption. A few architectural changes are needed to allow this to happen: The name nodes must use highly available shared storage to share the edit log. When a standby name node comes up, it reads up to the end of the shared edit log to synchronize its state with the active name node, and then continues to read new entries as they are written by the active name node. Data nodes must send block reports to both name nodes because the block mappings are stored in a name node's memory, and not on disk. Clients must be configured to handle name node failover, using a mechanism that is transparent to users. The secondary name node's role is subsumed by the standby, which takes periodic checkpoints of the active name node's namespace.

To understand the HDFS theory let's take an example. We have a file (file.txt) of size 200 MB. Manually we have taken a block of size 64 MB so the file.txt will be further sub divided into four sub files as follows:

- a.txt(64 MB),
- b.txt(64 MB),

c.txt(64 MB) and
d.txt(8 MB).

The advantage of HDFS is that the fourth block which is also of 64 MB but we are using only 8 MB of space so the rest free memory will not be wasted as in general file system rather it will be used for




Fig. 1 Storage process

further work. From the fig the client will send a request to name node that I want to store file.txt in your cluster. Now name node will take care of the request and create a meta data by taking the input about the file name, it's size and the blocks i.e the number of input split(a.txt, b.txt, c.txt and d.txt) and provide the data nodes number(1,3,5,7) as a response to request. Now the client is approaching datanode1 and store the a.txt file and so on. As the system is made up of commodity hardware so there is a chance of system failure and in case datanode1 becomes dead then the retrieval of data will not be possible and it will result in data lost when the client will require it. To overcome this problem HDFS has provided a solution, by default it will keep three replication of data. In this case node1 data is copied to data node 3 and 5 and a ack will be given back to data node1 that a.txt is stored and finally the data node1 will give the ack to the client that your file (a.txt) has been stored in data node.

B. MapReduce

MapReduce is a parallel programming model for writing distributed applications devised at Google for efficient processing of large amounts of data (multi-terabyte data-sets), on large clusters (thousands of nodes) of commodity. A MapReduce *job* is a unit of work that the client wants to be performed: it consists of the input data, the MapReduce program, and configuration information. Hadoop runs the job by dividing it into *tasks*, of which there are two types: map tasks and reduce tasks. The tasks are scheduled using YARN and run on nodes in the cluster. If a task fails, it will be automatically rescheduled to run on a different node.




Fig. 2 Retrieval of data

Now we will continue the HDFS example to see how mapping is done.

Client will send a job request in a form of programme (may be in java or in any scripting language) to job tracker which in turn will send a request to name node and ask where file.txt is?. The name node will check in file.txt in meta data and if it get the availability of the file.txt then it will reply the ack with meta data of file.txt. From that meta data it will get the data node where the files (block) is. Now the job tracker will send a request to task tracker, task tracker will check with the data node and send a ack back to job tracker. This process of sending request to data node through the task tracker and getting the +ve ack about the data is called mapping.

Combiner Functions: Many MapReduce jobs are limited by the bandwidth available on the cluster, so it pays to minimize the data transferred between map and reduce tasks. Hadoop allows the user to specify a *combiner function* to be run on the map output, and the combiner function's output forms the input to the reduce function. Because the combiner function is an optimization, Hadoop does not provide a guarantee of how many times it will call it for a particular map output record, if at all. In other words, calling the combiner function zero, one, or many times should produce the same output from the reducer.

IV. APPLICATION OF HADOOP

Big data applications solve and analyze real world problems using Hadoop and associated tools. Internet users and machine-to-machine connections are causing the data growth. Real time areas are defined following in which big data is used:

HEALTHCARE

Healthcare practices and policies differ tremendously around the world, there are three objectives regarding healthcare system [5]. The first objective is to improve the patient experience (including quality and satisfaction). Second, improving overall population health and reducing the cost of

health care and third is traditional methods have fallen short to manage healthcare and create modern technology to analyze large quantities of information. It is time consuming for clinical staff to Collecting massive amounts of data in healthcare. High-performance analytics are new technologies making easier to turn massive amounts of data into relevant and critical insights used to provide better care.

NETWORK SECURITY

Big data is a game changer in security technologies. The tremendous role of big data can be seen in network monitoring, forensics and SIEM [6]. Big data can also create a world where maintaining control over the revelation of our personal information is challenged constantly. Present analytical techniques don't work well at large scales and end up producing false positives that their efficacy is undermined and enterprises move to cloud architectures and gather much more data, the problem is becoming worse. Big data analytics is an effective solution for processing of large scale information as security is major concern in enterprises. Fraud detection is uses for big data analytics. Phone and credit card companies have conducted large-scale fraud detection for decades. Mainly big data tools are particularly suited to become fundamental for forensics and ATP.

MARKET AND BUSINESS

Big Data is the biggest game-changing opportunity for sales and marketing, since 20 years ago the Internet went main stream, because of the unprecedented array of insights into customer needs and behaviours it makes possible [7]. But many executives who agree that this is true aren't sure how to make the most of it and they also find themselves faced with overwhelming amounts of data and rapidly changing customer behaviours, organizational complexity and increased competitive pressures. According to Gartner, 50% internet connection between Internet of things (IoT) devices and number reached over 15 billion in 2011 and 30 billion by 2020[8]. Some companies are succeeding at turning that Big Data promise into reality. Those that use Big Data and analytics effectively show profitability and productivity rates that are 5–6% higher than those of their peers. The companies that succeed aren't the ones who have the most data, but the ones who use it best.

SPORTS

Sport, in business, an increasing volume of information is being collected and captured. Technological advances will fuel exponential growth in this area for the foreseeable future, as athletes are continuously monitored by tools as diverse as sports daily saliva, GPS systems and heart rate monitors tests. These statistics and many more like them are high performance in Big Data. These numbers there is a massive amount of potential insight and intelligence for trainers, administrators, coaches, athletes, sports medics and players. Statistics can be analyzed and collected to better understand what are the critical factors for optimum performance and success, in all facets of elite sport? Injury prevention,

competition, Preparation, and rehabilitation can all benefit by applying this approach. Recruitment, Scouting and retention can also be enhanced by these powerful principles. Keeping an eye on various information, a coach or a manager can easily and quickly understand which athletes and players need additional support, training, and guidance. Areas for reasons for success and improvement will be understood more clearly. Used consistently this is a powerful measure of progress and performance. [3]

EDUCATION SYSTEMS

By using big data analytics in field of education systems, remarkable results can be seen [8]. Data on students online behaviour can provide educators with important insights, such as if a student requires more attention, the class understanding of a topic is not clear, or if the course has to be modified. Students are required to answer accompanying questions as they go through the set of online content before class. By tracking the number of students that have completed the online module, the time taken and accuracy of their answers, a lecturer can be better informed of the profile of his students and modify the lesson plan accordingly. The analysis of data also clarify about the interest of student looking at time spent in online textbook, online lectures, notes etc. As result instructor can guide choosing the future path effectively.

V. FUTURE SCOPE

Microsoft Research's MyLifeBits project gives a glimpse of the archiving of personal information that may become common in the near future. MyLifeBits was an experiment where an individual's interactions like phone calls, emails, documents were captured electronically and stored for later access. The data gathered included a photo taken every minute, which resulted in an overall data volume of 1 gigabyte per month. When storage costs come down enough to make it feasible to store continuous audio and video, the data volume for a future MyLifeBits service will be many times that.

Machine logs, RFID readers, sensor networks, vehicle GPS traces, retail transactions etc. are generating vast amount of data in structured and unstructured form. Big data is able to process and store that data and probably in more amounts in near future. Hopefully, Hadoop will get better. New technologies and tools that have ability to record, monitor measure and combine all kinds of data around us, are going to be introduced soon. We will need new technologies and tools for data analysis, tracking and auditing information, sharing and managing, our own personal data in future. So many aspects of life health, education, telecommunication, marketing, sports and business etc. that manages big data world need to be polished in future.

VI. CONCLUSIONS

The ability to analyze and store massive amount of structured, unstructured and semi-structure data promises ongoing opportunities for academic institutes, businesses and

government organizations. However, a common horizontal big data analytics platform is necessary to support these varieties of real time applications that include healthcare, security, market and business, sports, education system, gaming industry, telecommunications and probably many others in future. The applications have been discussed in this paper. Furthermore, challenges of big data, 5 V's volume, velocity, variety, value, veracity and cloud enabled big data with models and types are also described in this paper. The main goal of our paper is to make a survey of various big data applications that are use in IT industries or organisation to store massive amount of data using technologies (Hadoop, HIVE, NoSQL, MapReduce and HPCC).

REFERENCES

- [1] These statistics were reported in a study entitled "The Digital Universe of Opportunities: Rich Data and the Increasing Value of the Internet of Things."
- [2] All figures are from 2013 or 2014. For more information, see Tom Groenfeldt, "At NYSE, The Data Deluge Overwhelms Traditional Databases"; Rich Miller, "Facebook Builds Exabyte Data Centers for Cold Storage"; Ancestry.com's "Company Facts"; Archive.org's "Petabox"; and the Worldwide LHC Computing Grid project's welcome page.
- [3] "Tom White", "The Definitive Guide 3rd Edition", O'Reilly, 12-14, 43- 50, May.2012.
- [4] "Hdfs architecture and its advantages" 14-15 <http://www.tutorialspoint.com/hadoop/index.htm>
- [5] Sam curry "big data fuels intelligence –driven security" RSA Security brief 2013
- [6] Sabia, Sheetal Kalra "Applications of big Data: Current Status and Future Scope" 5IEECSE31 Aug, Chandigarh http://www.irdindia.co.in/journal/journal_ijacte/pdf/vol3.pdf
- [7] "Big Data is the Future of Healthcare Bill Hamilton" cognizant 20-20 insights 2012
- [8] Anthony G. Picciano, "The Evolution of Big Data and Learning Analytics in American Higher Education" Journal of Asynchronous Learning Networks